

## Setup of 5-Armed Bandit

The bandit has 5 arms ({Arm[1],Arm[2],Arm[3],Arm[4],Arm[5]}),  
and is also supported with 2 contexts ({ArmContext[1],ArmContext[2]}).  
Each context follows standard normal distribution. The reward is given by:  
(Reward =  $i$  + ArmContext[1]\* $i$  – 0.1\*ArmContext[2]\* $i^2$  + RandN[])

When no context is available, the expected payoff of Arm[i] is i.  
Given the context, the expected payoff of Arm[i] is (i+ArmContext[1]\*i-0.1\*ArmContext[2]\*i^2).

```
In[301]:= RandN[] := RandomVariate[NormalDistribution[]];
ArmContext = <| 1 → RandN[], 2 → RandN[] |>;
ArmsList = {Arm[1], Arm[2], Arm[3], Arm[4], Arm[5]}

Out[303]:= {<| Reward → 9.52772, NoContextRegret → 4, WithContextRegret → 5.46842 |>,
<| Reward → 5.41839, NoContextRegret → 3, WithContextRegret → 4.86687 |>, <| Reward → 0.558041, NoContextRegret → 2, WithContextRegret → 0.166096 |>,
<| Reward → 8.06869, NoContextRegret → 1, WithContextRegret → 1.37113 |>, <| Reward → 7.49962, NoContextRegret → 0, WithContextRegret → 0. |> }

In[304]:= Arm[i_] := (*At each call to the arm[i] (i∈{1,2,3,4,5})),
do the following:*)
(Reward = i
+ ArmContext[1] * i
- 0.2 * ArmContext[2] * (6 - i)^2
+ RandN[];
(*This is the real reward returned by the arm.
The Algorithm can only observe this*)
NoContextRegret = 5
- i;
(*Regret = {Maximum expected reward arm}'s expected reward, here is 5
- this arm's expected reward*)
WithContextRegret = Max[Table[(ind + ArmContext[1] * ind - 0.1 * ArmContext[2] * ind^2), {ind, 5}]]
- (i + ArmContext[1] * i - 0.1 * ArmContext[2] * i^2);
(*Regret = {Maximum expected reward arm}'s expected reward
- this arm's expected reward*)
ArmContext = <| 1 → RandN[], 2 → RandN[] |>; (*Update a random context to display*)
<| "Reward" → Reward,
"NoContextRegret" → NoContextRegret,
"WithContextRegret" → WithContextRegret |>
(*Return the Actual reward and the Regret values*))
```

We allow for 1000 times of playing the bandit. The algorithm’s target is to maximize payoff / minimize regret.

```
In[383]:= TotalTimes = 1000
Out[383]:= 1000
```

## Bandit Algorithms Without Context

### Explore-First MAB Algorithm

This algorithm explores each arm by  $N$  times. For the rest, it choose to pull the arm with largest payoff given previous performance.

Mathematicians have proved that the optimal choice of  $N$  is:  $N = \left(\frac{T}{K}\right)^{2/3} \text{Log}[T]^{1/3}$  for  $K$ -armed bandits.

```
In[306]:= OptimalN = Round[ $\left(\frac{\text{TotalTimes}}{5}\right)^{2/3} \text{Log}[\text{TotalTimes}]^{1/3}$ ]
Out[306]:= 65
```

Implement the bandit algorithm:


```
In[307]:= ExploreFirstMABAlgorithm := (
RegretList = {};
RewardList = {};
(*Record regret & reward*)
MeanRewards = Table[Table[
Results = Arm[i];
AppendTo[RegretList, Results["NoContextRegret"]]; AppendTo[RewardList, Results["Reward"]];
Results["Reward"]
, {OptimalN}] // Mean, {i, 5}];
Print["Exploration Ended. Mean Rewards for 5 arms are:", MeanRewards];
(*This is the explored mean rewards from exploration phase. Print it out. *)
MaxArmIndex = MaximalBy[Range[5], MeanRewards[[#]] &][[1]];
Print["selected max-reward arm's index is:", MaxArmIndex];
(*This is the selected max-reward arm's index. *)
Table[
Results = Arm[MaxArmIndex];
AppendTo[RegretList, Results["NoContextRegret"]];
AppendTo[RewardList, Results["Reward"]];
, {TotalTimes - 5 * OptimalN}];
(*Choose the max-return arm in exploration phase. *)
)
```

```
In[308]:= ExploreFirstMABAlgorithm
```

Exploration Ended. Mean Rewards for 5 arms are:{0.409764, 1.8188, 2.31432, 4.12468, 4.28383}

selected max-reward arm's index is:5

```
In[312]:= ListPlot[ {RegretList, RewardList}, PlotLabel → "Regret/Reward over time",
  PlotLegends → {"Regret over time", "Reward over time"}, PlotStyle → Directive[{PointSize[Small]}] ]
```



```
In[313]:= ListLinePlot[ {Accumulate[RegretList], Accumulate[RewardList]},
  PlotLabel → "Total Regret/Reward over time", PlotLegends → {"Total Regret over time", "Total Reward over time"} ]
```



```
In[314]:= Total[RewardList] / TotalTimes
```

Out[314]= 4.2679

## ϵ-Greedy MAB Algorithm

Mathematicians have proved that, to minimize the upper-bound of regret, the optimal choice of  $\epsilon$  is:  $\epsilon = (T)^{-1/3} (5 \text{ Log}[T])^{1/3}$  for K-armed bandits.

```
In[315]:= Optimalϵ = N[ (TotalTimes)^{-1/3} (5 Log[TotalTimes])^{1/3}]
```


Out[315]= 0.325663

Implement the bandit algorithm:

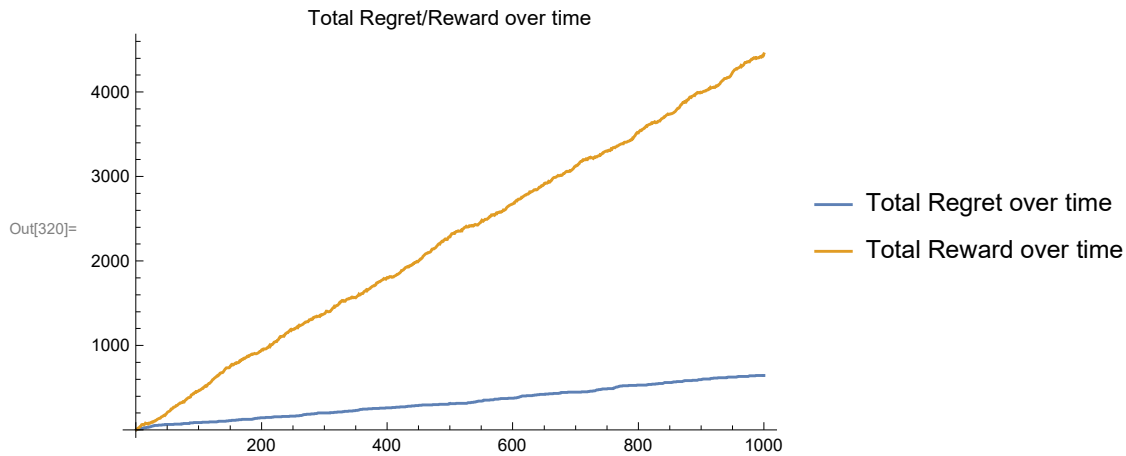
```
In[316]:= eGreedyMABAlgorithm := (
  RegretList = {};
  RewardList = {};
  (*Record regret & reward*)
  ExpectedRewards = {10, 10, 10, 10, 10};
  ArmDrawCounts = {0, 0, 0, 0, 0};
  (*initial expectation.*)
  Table[
    If[RandomReal[] < Optimalϵ,
      (*flip a coin*)
      tmpArm = RandomInteger[{1, 5}];
      Results = Arm[tmpArm];
      ExpectedRewards[[tmpArm]] = (ExpectedRewards[[tmpArm]] * ArmDrawCounts[[tmpArm]] + Results["Reward"]) / (ArmDrawCounts[[tmpArm]] + 1);
      ArmDrawCounts[[tmpArm]]++;
      AppendTo[RegretList, Results["NoContextRegret"]];
      AppendTo[RewardList, Results["Reward"]];,
      (*randomly draw an arm*)
      Results = Arm[MaxArmIndex];
      ExpectedRewards[[MaxArmIndex]] =
        (ExpectedRewards[[MaxArmIndex]] * ArmDrawCounts[[MaxArmIndex]] + Results["Reward"]) / (ArmDrawCounts[[MaxArmIndex]] + 1);
      ArmDrawCounts[[MaxArmIndex]]++;
      AppendTo[RegretList, Results["NoContextRegret"]];
      AppendTo[RewardList, Results["Reward"]];
      (*else, just draw the best arm so far*)
    ];
    MaxArmIndex = MaximalBy[Range[5], ExpectedRewards[[#]] &][[1]];
    (*update maxarmindex*)
    , {TotalTimes}];
  (*Choose the max-return arm in exploration phase.*)
)
```

```
In[317]:= eGreedyMABAlgorithm
```

```
In[319]:= ListPlot[ {RegretList, RewardList}, PlotLabel → "Regret/Reward over time",
  PlotLegends → {"Regret over time", "Reward over time"}, PlotStyle → Directive[{PointSize[Small]}] ]
```



```
In[320]:= ListLinePlot[ {Accumulate[RegretList], Accumulate[RewardList]},
  PlotLabel → "Total Regret/Reward over time", PlotLegends → {"Total Regret over time", "Total Reward over time"} ]
```



```
In[321]:= Total[RewardList] / TotalTimes
```

Out[321]= 4.44966

## Thompson-Sampling MAB Algorithm

Thompson sampling algorithm explores each arm at the Bayesian posterior probability of the arm being max-return arm. The algorithm requires a “prior” distribution, so the calculation is pretty difficult.

When no contexts are available, the distribution of each arm’s payoff is dependent on i:

```
In[105]:= TransformedDistribution[
  i + i * u - 0.1 v * i^2 + w,
  {u ≈ NormalDistribution[], v ≈ NormalDistribution[], w ≈ NormalDistribution[]}]
```

Out[105]= NormalDistribution[0. + i,  $\sqrt{1 + i^2 + 0.01 i^4}$ ]

Typically, we don’t have a prior knowledge so accurate that we know exactly how the returns are distributed. Instead we usually have to infer from the observed data. Mathematica provides built-in function to do this task: EstimatedDistribution.

We specify Normal distribution as a prior distribution.

```
In[ ]:= FindDistributionParameters[data, NormalDistribution[μ, θ]]
```

And, as we inferred the distribution parameter, from 计量经济学知识, we know the error of the estimated parameters. Thus, we can calculate the probability of one arm surpassing the other.

```
In[70]:= Probability[x ≤ y, {x ≈ NormalDistribution[μx,  $\frac{\theta_x}{\sqrt{\text{Len}_{xdata} - 1}}$ ], y ≈ NormalDistribution[μy,  $\frac{\theta_y}{\sqrt{\text{Len}_{ydata} - 1}}$ ]}]
```

Out[70]=  $\frac{1}{2} \text{Erfc}\left[\frac{\mu_x - \mu_y}{\sqrt{2} \sqrt{\frac{\theta_x^2}{-1 + \text{Len}_{xdata}} + \frac{\theta_y^2}{-1 + \text{Len}_{ydata}}}}\right]$

```
In[ ]:= FindDistributionParameters[ArmDrawResults[[1]], NormalDistribution[μ, θ]]
```

The probability of certain arm being the highest-return arm, is the product of expressions like this.


Also, we want to give the Thomson algorithm a little bit of prior knowledge before observing. We will spare 20 observing chances for each arm to provide a little bit prior knowledge to infer from.

Implement the bandit algorithm:


```
In[322]:= ThompsonSamplingMABAlgorithm := (
  RegretList = {};
  RewardList = {};
  (*Record regret & reward*)
  ArmDrawResults = Table[Table[
    Results = Arm[i];
    AppendTo[RegretList, Results["NoContextRegret"]]; AppendTo[RewardList, Results["Reward"]];
    Results["Reward"] + 0(*optimistic boost. helps the algorithm to explore better*)
    , {20}], {i, 5}];
  DistParams = Table[({μ, θ} /. FindDistributionParameters[ArmDrawResults[[i]], NormalDistribution[μ, θ]]), {i, 5}];
  (*μ, θ*)
  ProbabilityOfBeingMax = Table[
    Product[ $\frac{1}{2}$  Erfc[(DistParams[[j, 1]] - DistParams[[ind, 1]]) /  $\sqrt{2}$   $\sqrt{(\text{DistParams}[[j, 2]]^2 / (-1 + \text{Length}[\text{ArmDrawResults}[[j]]) + \text{DistParams}[[ind, 2]]^2 / (-1 + \text{Length}[\text{ArmDrawResults}[[ind]])}$ )]], {j, Drop[Range[5], {ind}]}]
    , {ind, 5}];
  (*get initial startup knowledge, plus give every arm a optimistic boost.*)
  Table[
    tmpArm = RandomChoice[ProbabilityOfBeingMax → Range[5]];
    (*this is the sampling algorithm's choice*)
    Results = Arm[tmpArm];
    AppendTo[RegretList, Results["NoContextRegret"]];
    AppendTo[RewardList, Results["Reward"]];
    (*record regrets & rewards*)
    AppendTo[ArmDrawResults[[tmpArm]], Results["Reward"]];
    (*the Algorithm record the results.*)
    DistParams[[tmpArm]] = ({μ, θ} /. FindDistributionParameters[ArmDrawResults[[tmpArm]], NormalDistribution[μ, θ]]);
    (*update the params of distribution*)
    ProbabilityOfBeingMax = Table[
      Product[ $\frac{1}{2}$  Erfc[(DistParams[[j, 1]] - DistParams[[ind, 1]]) /  $\sqrt{2}$   $\sqrt{(\text{DistParams}[[j, 2]]^2 / (-1 + \text{Length}[\text{ArmDrawResults}[[j]]) + \text{DistParams}[[ind, 2]]^2 / (-1 + \text{Length}[\text{ArmDrawResults}[[ind]])}$ )]], {j, Drop[Range[5], {ind}]}]
      , {ind, 5}];
      (*update bayesian probability*)
      , {TotalTimes - 5 * 20}];
      (*Choose the max-return arm in exploration phase.*)
    ]
)
```

```
In[323]:= ThompsonSamplingMABAlgorithm
```

```
In[324]:= ListPlot[ {RegretList, RewardList}, PlotLabel → "Regret/Reward over time",
  PlotLegends → {"Regret over time", "Reward over time"}, PlotStyle → Directive[{PointSize[Small]}] ]
```



```
In[325]:= ListLinePlot[ {Accumulate[RegretList], Accumulate[RewardList]},
  PlotLabel → "Total Regret/Reward over time", PlotLegends → {"Total Regret over time", "Total Reward over time"} ]
```



Note that TS algorithm is pretty unstable. The initial draws will affect a lot. Sometimes it almost always choose the right arm, while sometime not that lucky.

```
In[326]:= Total[RewardList] / TotalTimes
```

```
Out[326]:= 4.66514
```

UCB MAB Algorithm


A final MAB algorithm for non-contextual algorithm is UCB algorithm. This algorithm is also adaptive, and it always chooses the arm with maximum Upper Confidence Bound.

Implement the bandit algorithm:


```
In[338]:= UCBMABAlgorithm := (
  RegretList = {};
  RewardList = {};
  (*Record regret & reward*)
  ArmDrawResults = {{20}, {20}, {20}, {20}, {20}};
  UCBList = {500, 500, 500, 500, 500};
  (*have to make an optimistic projection*)
  Table[
    tmpArm = MaximalBy[Range[5], UCBList[[#]] &][[1]];
    (*this is the sampling algorithm's choice*)
    Results = Arm[tmpArm];
    AppendTo[RegretList, Results["NoContextRegret"]];
    AppendTo[RewardList, Results["Reward"]];
    (*record regrets & rewards*)
    AppendTo[ArmDrawResults[[tmpArm]], Results["Reward"]];
    (*the Algorithm record the results.*)
    UCBList[[tmpArm]] = Quantile[EstimatedDistribution[ArmDrawResults[[tmpArm]], NormalDistribution[μ, θ]], 1 - 0.05];
    (*update UCB. we set the confidence level at 0.05 (one-sided)*)
    , {TotalTimes}];
)
```

```
In[339]:= UCBMABAlgorithm
```

```
In[340]:= ListPlot[ {RegretList, RewardList}, PlotLabel → "Regret/Reward over time",
  PlotLegends → {"Regret over time", "Reward over time"}, PlotStyle → Directive[{PointSize[Small]}] ]
```



```
In[341]:= ListLinePlot[ {Accumulate[RegretList], Accumulate[RewardList]},
  PlotLabel → "Total Regret/Reward over time", PlotLegends → {"Total Regret over time", "Total Reward over time"} ]
```



```
In[342]:= Total[RewardList] / TotalTimes
```

```
Out[342]:= 4.5904
```

UCB is even more unstable than TS, and its performance under large-variance scenario is just poor. Basically UCB forbade itself from exploring arms from its maximizing greedy behaviour.

## Contextual Bandit Algorithms

### Contextual Thompson-Sampling MAB Algorithm

Extension to naive Thompson Sampling algorithm. Define a “prior” distribution with consideration of contexts.

We specify Normal distribution with its mean dependent on context as a prior distribution. Effectively, the context-generated distribution will also be a normal distribution.

```
In[*]:= FindDistributionParameters[data, NormalDistribution[μ, θ]]
```

Calculate the probability of one arm surpassing the other (Same as before).

$$\text{In[*]} := \text{Probability}\left[\mathbf{x} \leq \mathbf{y}, \left\{\mathbf{x} \approx \text{NormalDistribution}\left[\mu_x, \frac{\theta_x}{\sqrt{\text{Len}_{xdata} - 1}}\right], \mathbf{y} \approx \text{NormalDistribution}\left[\mu_y, \frac{\theta_y}{\sqrt{\text{Len}_{ydata} - 1}}\right]\right\}\right]$$

$$\text{Out[*]} := \frac{1}{2} \text{Erfc}\left[\frac{\mu_x - \mu_y}{\sqrt{2} \sqrt{\frac{\theta_x^2}{-1 + \text{Len}_{xdata}} + \frac{\theta_y^2}{-1 + \text{Len}_{ydata}}}}\right]$$

The probability of certain arm being the highest-return arm, is the product of expressions like this. Note we have to estimate the parameters using nonlinear model fit.

Also, we give the Thomson algorithm a little bit of optimistic prior knowledge before observing.


Implement the bandit algorithm:

```
In[377]:= ThompsonSamplingContextualMABAlgorithm := (
  RegretList = {};
  RewardList = {};
  (*Record regret & reward*)
  ArmDrawResults = {{ {1, 0, 21}, {0, 1, 10}, {-1, 0, 19.5}, {0, -0.5, 10} },
    { {1, 0, 21}, {0, 1, 10}, {-1, 0, 19.5}, {0, -0.5, 10} }, { {1, 0, 21}, {0, 1, 10}, {-1, 0, 19.5}, {0, -0.5, 10} },
    { {1, 0, 21}, {0, 1, 10}, {-1, 0, 19.5}, {0, -0.5, 10} }, { {1, 0, 21}, {0, 1, 10}, {-1, 0, 19.5}, {0, -0.5, 10} } };
  (*optimistic prior expectation*)
  Table[
    DistParams = Table[
      nlm = NonlinearModelFit[ArmDrawResults[[ind]], a + b x + c y, {a, b, c}, {x, y}];
      bestFitParams = ({a, b, c} /. nlm["BestFitParameters"]);
      bestFitParamsErrors = nlm["ParameterErrors"];
      EstVariance = nlm["EstimatedVariance"];
      tmpDist = (TransformedDistribution[
        a + b x + c y + ε, {ε ≈ NormalDistribution[0, √EstVariance], a ≈ NormalDistribution[bestFitParams[[1]], bestFitParamsErrors[[1]]],
        b ≈ NormalDistribution[bestFitParams[[2]], bestFitParamsErrors[[2]]],
        c ≈ NormalDistribution[bestFitParams[[3]], bestFitParamsErrors[[3]]]}
      ] /. {x → ArmContext[1], y → ArmContext[2]});
      {tmpDist[[1]], tmpDist[[2]]}
    , {ind, 5}];
  ProbabilityOfBeingMax = Table[
    Product[
      1/2 Erfc[(DistParams[[j, 1]] - DistParams[[ind, 1]]) / (√2 √(DistParams[[j, 2]]^2 / (-1 + Length[ArmDrawResults[[j]])) +
        DistParams[[ind, 2]]^2 / (-1 + Length[ArmDrawResults[[ind]])))] , {j, Drop[Range[5], {ind}]}]
    , {ind, 5}];
  (*find the Distribution Parameters, and the probability*)
  tmpArm = RandomChoice[ProbabilityOfBeingMax → Range[5]];
  (*this is the sampling algorithm's choice*)
  Results = Arm[tmpArm];
  AppendTo[RegretList, Results["WithContextRegret"]];
  AppendTo[RewardList, Results["Reward"]];
  (*record regrets & rewards*)
  AppendTo[ArmDrawResults[[tmpArm]], {ArmContext[1], ArmContext[2], Results["Reward"]}];
  (*the Algorithm record the context and results.*)
  , {TotalTimes}];
)
```

```
In[392]:= ThompsonSamplingContextualMABAlgorithm
```



```
In[393]:= ListPlot[ {RegretList, RewardList}, PlotLabel → "Regret/Reward over time",
  PlotLegends → {"Contextual Regret over time", "Reward over time"}, PlotStyle → Directive[{PointSize[Small]}] ]
```



```
In[394]:= ListLinePlot[ {Accumulate[RegretList], Accumulate[RewardList]}, PlotLabel → "Total Regret/Reward over time",
  PlotLegends → {"Total Contextual Regret over time", "Total Reward over time"} ]
```



```
In[395]:= Total[RewardList] / TotalTimes
```

Out[395]= 4.48149

## Contextual UCB: LinUCB Algorithm

UCB algorithm can also be modified to utilize context information.


Implement the bandit algorithm:

```
In[483]:= LinUCBContextualMABAlgorithm := (
  RegretList = {};
  RewardList = {};
  (*Record regret & reward*)
  ArmDrawResults = {{ {1, 0, 15}, {0, 1, 10}, {-1, 0, 12.5}, {0, -0.5, 10} },
    { {1, 0, 15}, {0, 1, 10}, {-1, 0, 12.5}, {0, -0.5, 10} }, { {1, 0, 15}, {0, 1, 10}, {-1, 0, 12.5}, {0, -0.5, 10} },
    { {1, 0, 15}, {0, 1, 10}, {-1, 0, 12.5}, {0, -0.5, 10} }, { {1, 0, 15}, {0, 1, 10}, {-1, 0, 12.5}, {0, -0.5, 10} } };
  UCBLList = {10000, 10000, 10000, 10000, 10000};
  (*have to make an optimistic projection*)
  Table[
    UCBLList = UCBLList * 0.5 + Table[
      nlm = NonlinearModelFit[ArmDrawResults[[ind]], a + b x + c y, {a, b, c}, {x, y}];
      bestFitParams = ({a, b, c} /. nlm["BestFitParameters"]);
      bestFitParamsErrors = nlm["ParameterErrors"];
      EstVariance = nlm["EstimatedVariance"];
      Quantile[TransformedDistribution[
        a + b x + c y + e, {e ≈ NormalDistribution[0, sqrt[EstVariance]], a ≈ NormalDistribution[bestFitParams[[1]], bestFitParamsErrors[[1]]],
        b ≈ NormalDistribution[bestFitParams[[2]], bestFitParamsErrors[[2]]],
        c ≈ NormalDistribution[bestFitParams[[3]], bestFitParamsErrors[[3]]]}
      ], 1 - 0.00005] /. {x → ArmContext[1], y → ArmContext[2]}
    , {ind, 5}];

  (*find the UCB FOR THE CONTEXT STATUS QUO. we set the confidence level at 0.00005 (one-sided)*)
  tmpArm = MaximalBy[Range[5], UCBLList[[#]] &][[1]];
  (*this is the sampling algorithm's choice*)
  Results = Arm[tmpArm];
  AppendTo[RegretList, Results["WithContextRegret"]];
  AppendTo[RewardList, Results["Reward"]];
  (*record regrets & rewards*)
  AppendTo[ArmDrawResults[[tmpArm]], {ArmContext[1], ArmContext[2], Results["Reward"]}];
  (*the Algorithm record the context and results.*)
  , {TotalTimes}];
)
```

In[492]:= LinUCBContextualMABAlgorithm

```
In[493]:= ListPlot[ {RegretList, RewardList}, PlotLabel → "Regret/Reward over time",  
PlotLegends → {"Contextual Regret over time", "Reward over time"}, PlotStyle → Directive[{PointSize[Small]}] ]
```



```
In[494]:= ListLinePlot[ {Accumulate[RegretList], Accumulate[RewardList]}, PlotLabel → "Total Regret/Reward over time",  
PlotLegends → {"Total Contextual Regret over time", "Total Reward over time"} ]
```



```
In[495]:= Total[RewardList] / TotalTimes  
Out[495]= 4.92722
```

Contextual bandits have greater potential to optimize reward, because they observe better information. You can see contextual regret is still pretty high, but the performance has also exceeded non-contextual bandits.